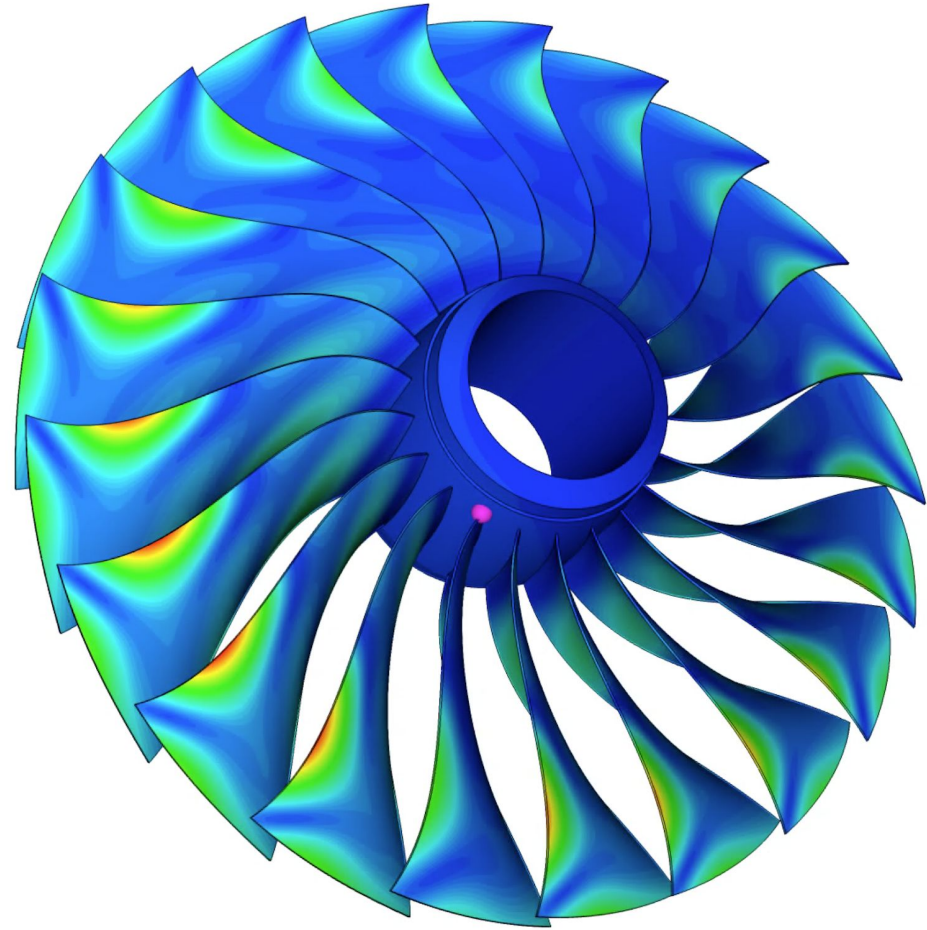


GPU accelerated computing for Finite Element Method




Introduction to GPU

GPU programming - Basics

Thread

- A thread is a single sequential flow of control within a program.
- A sequential code in one processor has one thread.



Introduction to GPU

GPU programming - C library functions

Pointer

- A variable that points to the storage/memory address of another variable.
 - A variable of type certain type will store a value

```
int v = 0;
```

- This variable has its address (where it is located the memory). This address can be obtained by using '&'

&v

- A pointer stores the address of the variable

```
int *y = &v;
```

- The value of the variable can be accessed using the variable or the pointer

v
*y

Introduction to GPU

GPU programming - C library functions


malloc()

- Dynamically allocates a single large block of memory
 - Syntax

```
pointer = (type*) malloc(byte size)
```

- Example

```
n = 5;  
int *p;  
p = (int*)malloc(n * sizeof(int));
```



Introduction to GPU

GPU programming syntax

EXERCISE 0 : Vector Addition using C program

Source code

https://github.com/sivasanarul/FEMwithGPU/tree/master/EX1_vector_addition

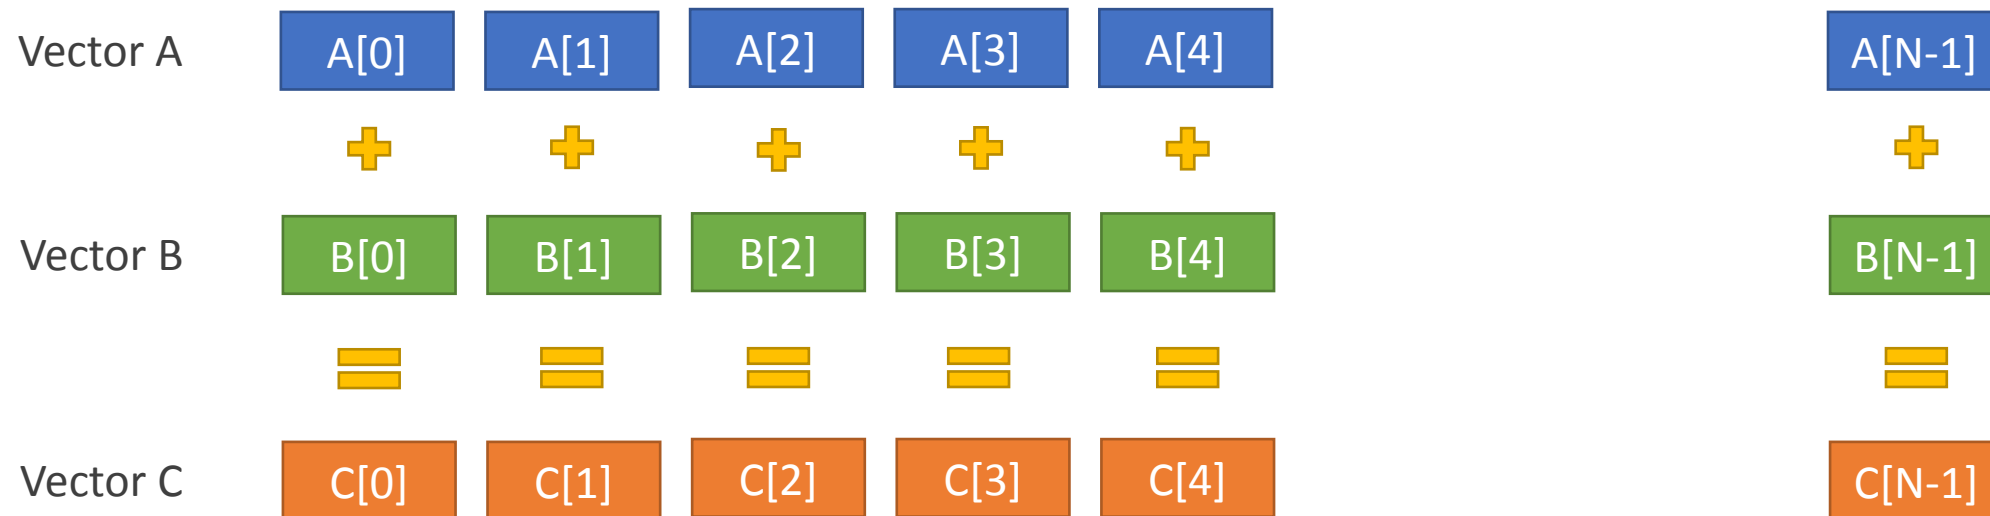
vector_add.c

Introduction to GPU

GPU programming

Vector Addition : `vector_addition()`

$$A + B = C$$



Introduction to GPU

GPU programming - C program for vector addition

C programming – Vector addition

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #define N 10000000
4
5  int main(){
6      float *a, *b, *c;
7
8      a = (float*)malloc(sizeof(float) * N);
9      b = (float*)malloc(sizeof(float) * N);
10     c = (float*)malloc(sizeof(float) * N);
11
12     // Initialize array
13     for(int i = 0; i < N; i++){
14         a[i] = 1.0f; b[i] = 2.0f;
15     }
16
17
18
19     clock_t t;
20     t = clock();
21     // Main function
22     for(int i=0;i<N;i++){
23         c[i] = a[i] + b[i];}
24     t = clock() - t;
25     double time_taken = ((double)t)/CLOCKS_PER_SEC; // in seconds
26     printf("fun() took %f seconds to execute \n", time_taken);
27
28     free(a); free(b); free(c);
29
30 }
```

```
6      float *a, *b, *c;
7
8      a = (float*)malloc(sizeof(float) * N);
9      b = (float*)malloc(sizeof(float) * N);
10     c = (float*)malloc(sizeof(float) * N);
```

Memory Allocation
for the variables



```
12     // Initialize array
13     for(int i = 0; i < N; i++){
14         a[i] = 1.0f; b[i] = 2.0f;
15     }
```

Initializing the
variables



```
21     // Main function
22     for(int i=0;i<N;i++){
23         c[i] = a[i] + b[i];}
```

Addition of vectors



```
28     free(a); free(b); free(c);
```

Deallocation of
Memory

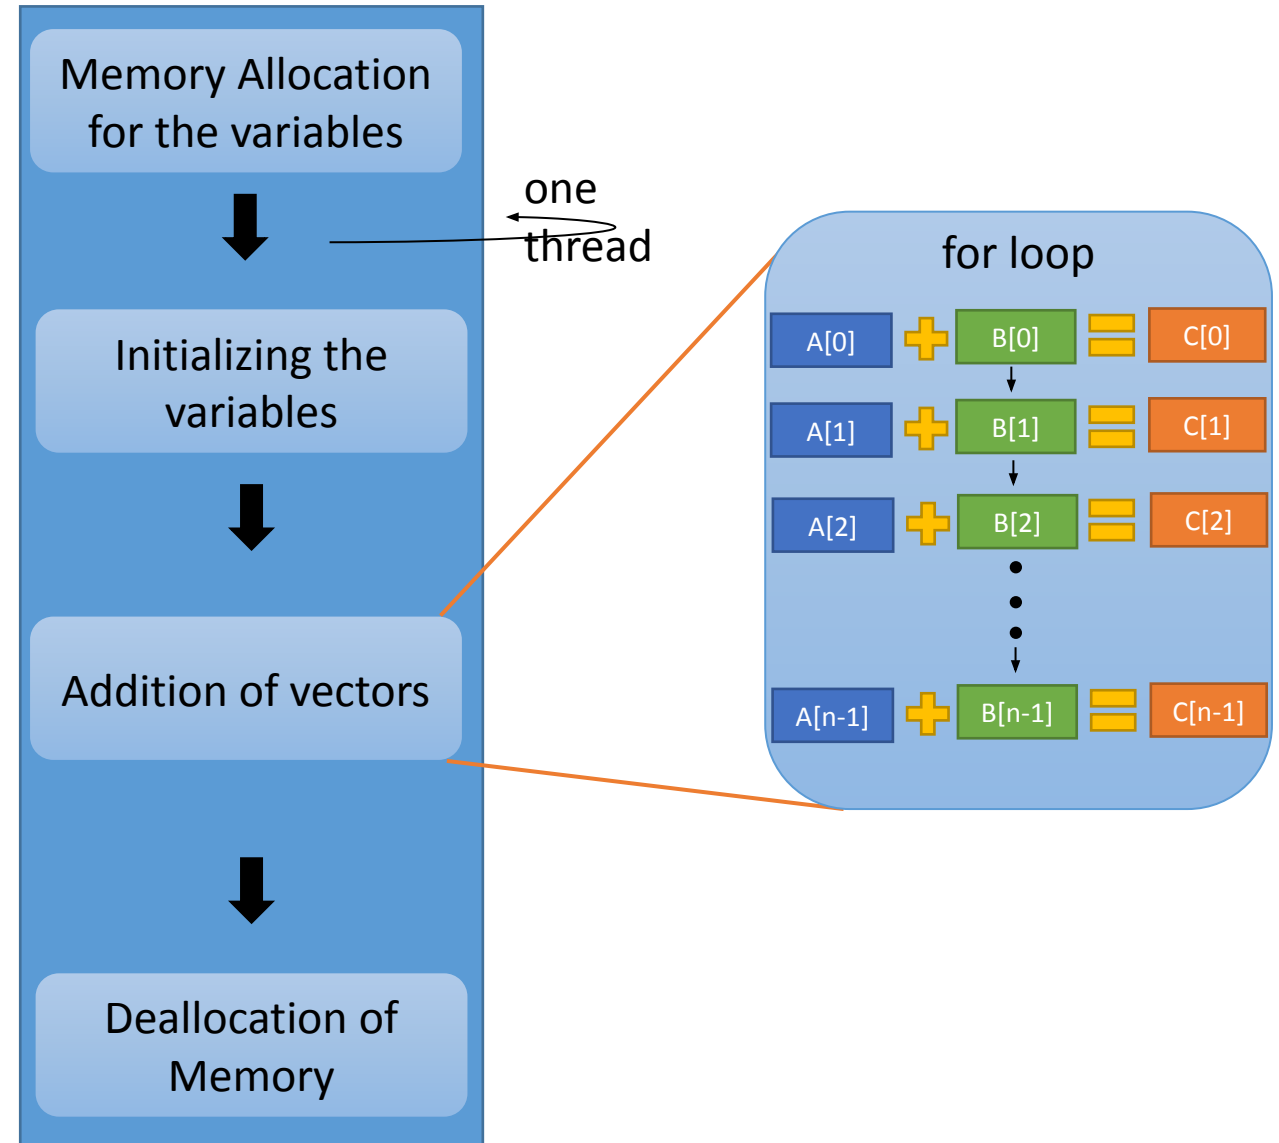
Introduction to GPU

GPU programming - C program for vector addition

C programming – Vector addition

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #define N 100000000
4
5  int main(){
6      float *a, *b, *c;
7
8      a = (float*)malloc(sizeof(float) * N);
9      b = (float*)malloc(sizeof(float) * N);
10     c = (float*)malloc(sizeof(float) * N);
11
12     // Initialize array
13     for(int i = 0; i < N; i++){
14         a[i] = 1.0f; b[i] = 2.0f;
15     }
16
17
18
19     clock_t t;
20     t = clock();
21     // Main function
22     for(int i=0;i<N;i++){
23         c[i] = a[i] + b[i];
24     }
25     t = clock() - t;
26     double time_taken = ((double)t)/CLOCKS_PER_SEC; // in seconds
27     printf("fun() took %f seconds to execute \n", time_taken);
28
29     free(a); free(b); free(c);
30 }
```

Serial
Code

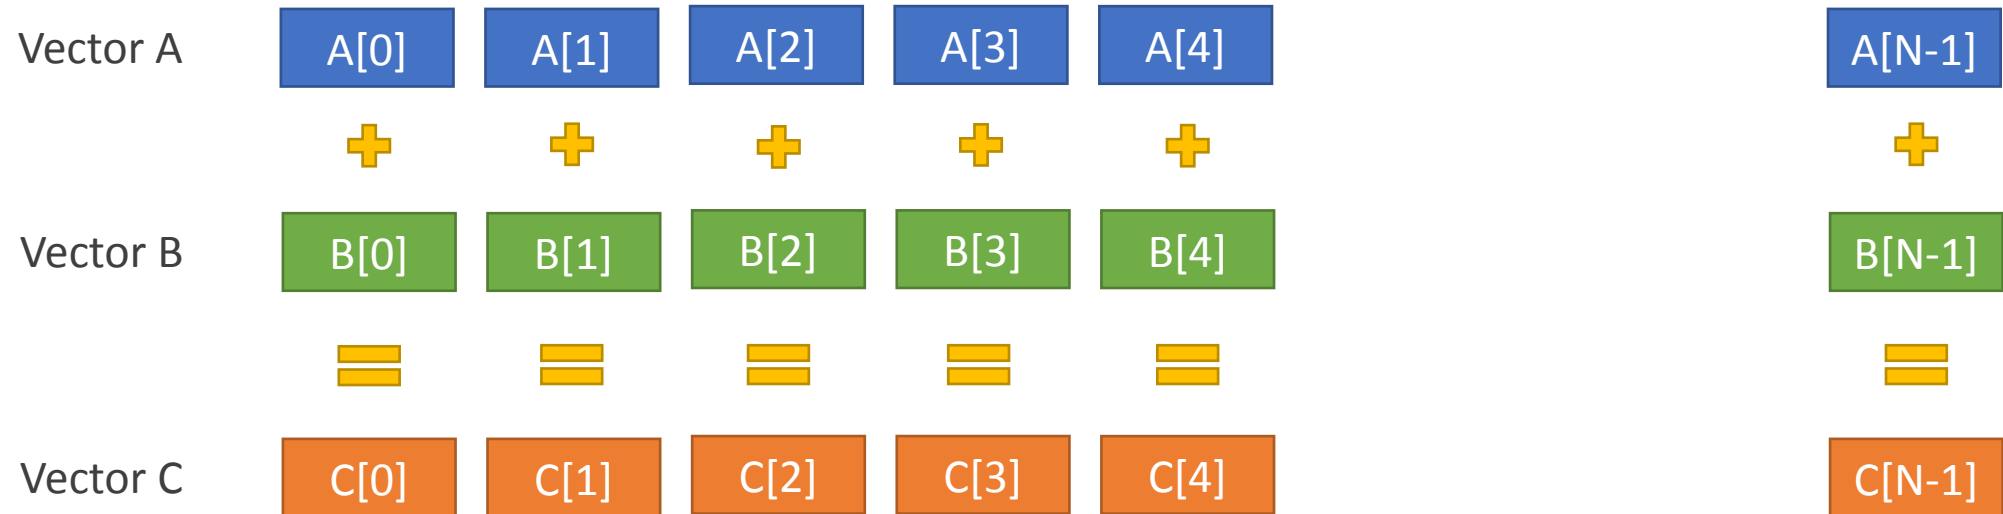


Introduction to GPU

GPU programming

Vector Addition : `vector_addition()`

$$A + B = C$$



- Divide array into blocks cuda
- threads into multiple blocks

Introduction to GPU

GPU programming - C program for vector addition, heterogenous program

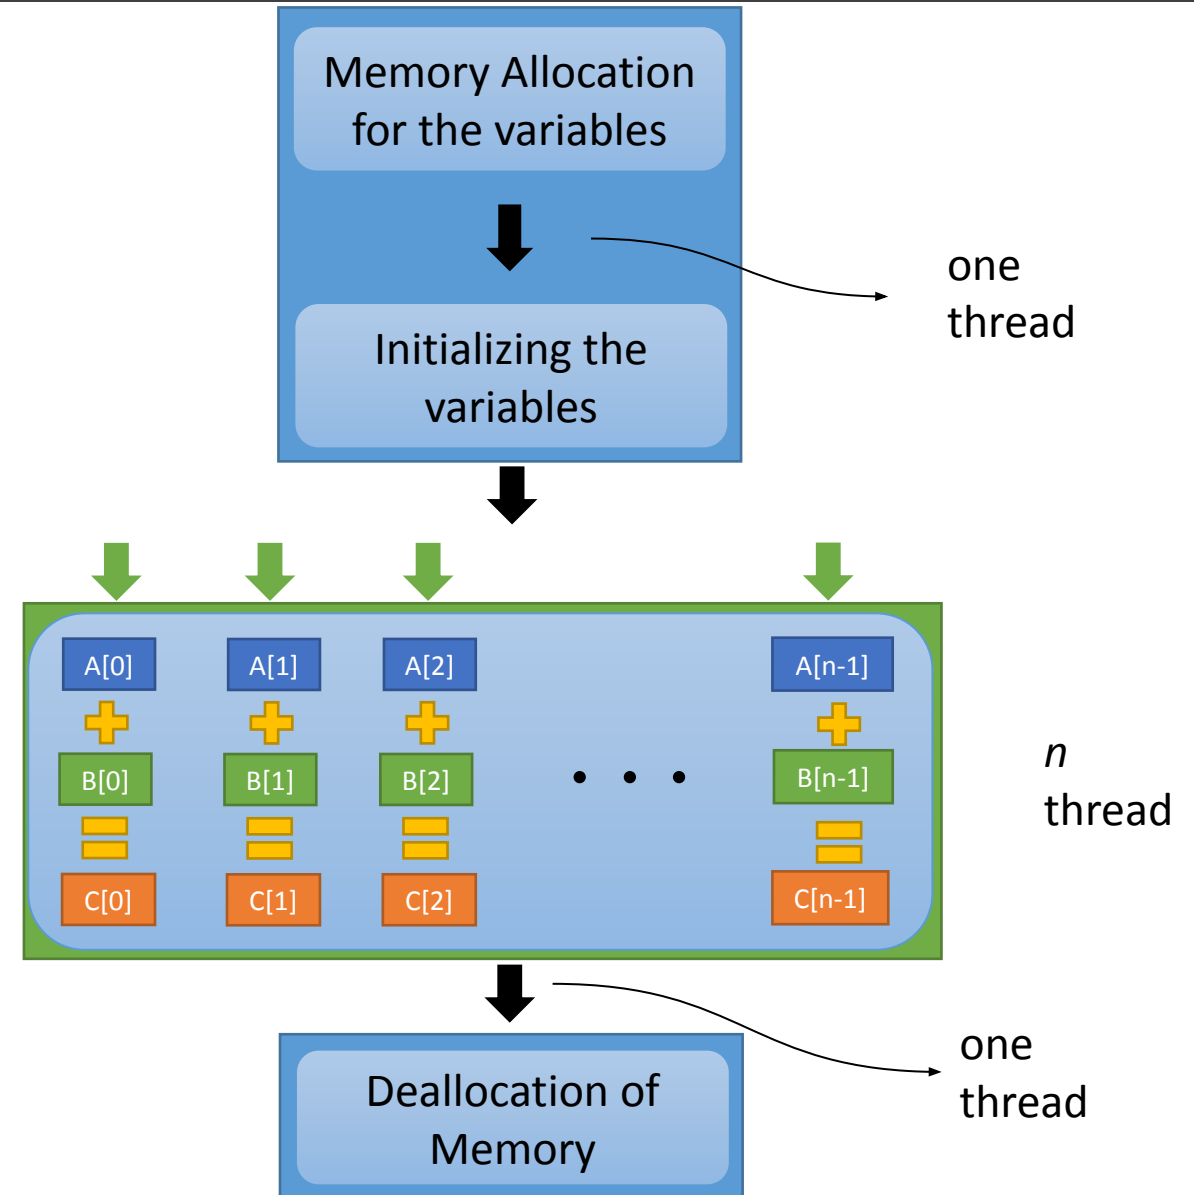
CUDA programming Model – Heterogenous Computing

```
16 int main(){
17     float *a, *b, *out;
18     float *d_a, *d_b, *d_out;
19
20     a = (float*)malloc(sizeof(float) * N);
21     b = (float*)malloc(sizeof(float) * N);
22     out = (float*)malloc(sizeof(float) * N);
23
24     // Initialize array
25     for(int i = 0; i < N; i++){
26         a[i] = 1.0f;
27         b[i] = 2.0f;
28     }
29
30     // Allocate device memory for a
31     cudaMalloc((void**)&d_a, sizeof(float)*N);
32     cudaMalloc((void**)&d_b, sizeof(float)*N);
33     cudaMalloc((void**)&d_out, sizeof(float)*N);
34
35     // Transfer data from host to device memory
36     cudaMemcpy(d_a, a, sizeof(float)*N, cudaMemcpyHostToDevice);
37     cudaMemcpy(d_b, b, sizeof(float)*N, cudaMemcpyHostToDevice);
38
39     // Main function
40     int block_size = 256;
41     int grid_size = (N+block_size)/block_size;
42     vector_add<<<grid_size, block_size>>>>(d_out, d_a, d_b, N);
43
44     cudaMemcpy(out, d_out, sizeof(float)*N, cudaMemcpyDeviceToHost);
45
46     // Deallocate device memory
47     cudaFree(d_a);
48     cudaFree(d_b);
49     cudaFree(d_out);
50
51     // Deallocate host memory
52     free(a);
53     free(b);
54     free(out);
55 }
56 }
```

Serial Code

Parallel Code

Serial Code



Introduction to GPU

GPU programming - Heterogenous program


CUDA programming Model – Heterogenous Computing

```
16 int main(){
17     float *a, *b, *out;
18     float *d_a, *d_b, *d_out;
19
20     a = (float*)malloc(sizeof(float) * N);
21     b = (float*)malloc(sizeof(float) * N);
22     out = (float*)malloc(sizeof(float) * N);
23
24     // Initialize array
25     for(int i = 0; i < N; i++){
26         a[i] = 1.0f;
27         b[i] = 2.0f;
28     }
29
30     // Allocate device memory for a
31     cudaMalloc((void**)&d_a, sizeof(float)*N);
32     cudaMalloc((void**)&d_b, sizeof(float)*N);
33     cudaMalloc((void**)&d_out, sizeof(float)*N);
34
35     // Transfer data from host to device memory
36     cudaMemcpy(d_a, a, sizeof(float)*N, cudaMemcpyHostToDevice);
37     cudaMemcpy(d_b, b, sizeof(float)*N, cudaMemcpyHostToDevice);
38
39     // Main function
40     int block_size = 256;
41     int grid_size = (N+block_size)/block_size;
42     vector_add<<<grid_size, block_size>>>>(d_out, d_a, d_b, N);
43
44     cudaMemcpy(out, d_out, sizeof(float)*N, cudaMemcpyDeviceToHost);
45
46     // Deallocate device memory
47     cudaFree(d_a);
48     cudaFree(d_b);
49     cudaFree(d_out);
50
51     // Deallocate host memory
52     free(a);
53     free(b);
54     free(out);
55
56 }
```

Serial Code

Parallel Code

Serial Code



Serial Code
in host

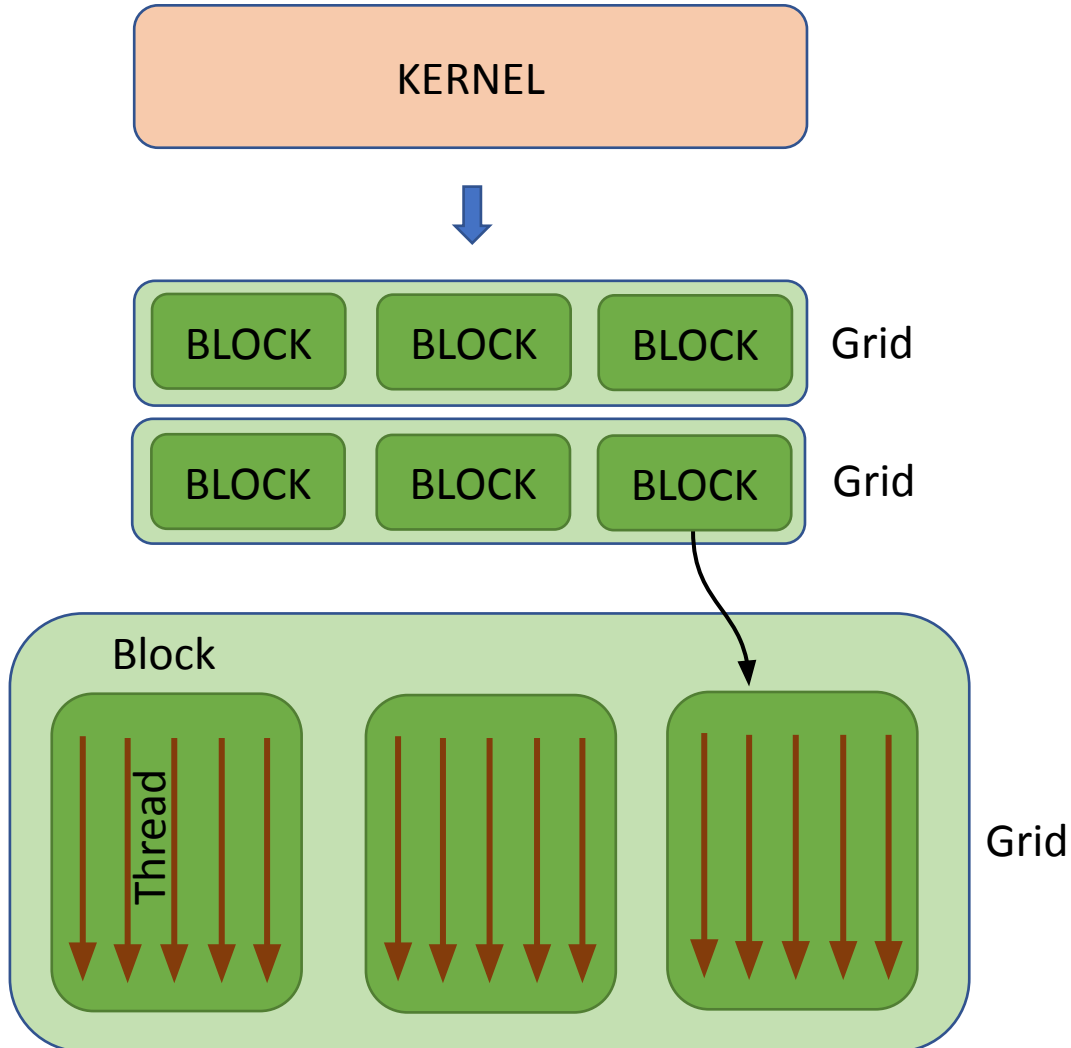
Parallel Code
in device

Serial Code
in host

Introduction to GPU

Introduction of GPU


A kernel is a function that is run in a GPU



- A kernel is split into grids
- A grid is split into blocks
- Each block is collection of threads

Introduction to GPU


Introduction of GPU



cores



Streaming Multiprocessor –
collection of cores





CUDA-enabled GPU –
Collection of Streaming
Multiprocessor

Introduction to GPU

Introduction of GPU

A kernel is a function that is run in a GPU



A kernel is executed in a CUDA-enabled GPU

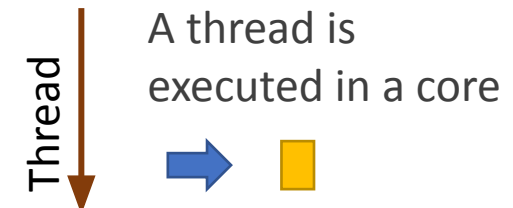


Multiple kernels can be executed in one CUDA-enabled GPU

A block is executed in a Streaming Multiprocessor



Multiple blocks can be executed in one Streaming multiprocessor



Introduction to GPU

Introduction of GPU programming


Vector Addition : `vector_addition()`

- ❑ In one core as one thread
- ❑ In one streaming multiprocessor as one block
- ❑ In the entire GPU device as multiple blocks

Introduction to GPU

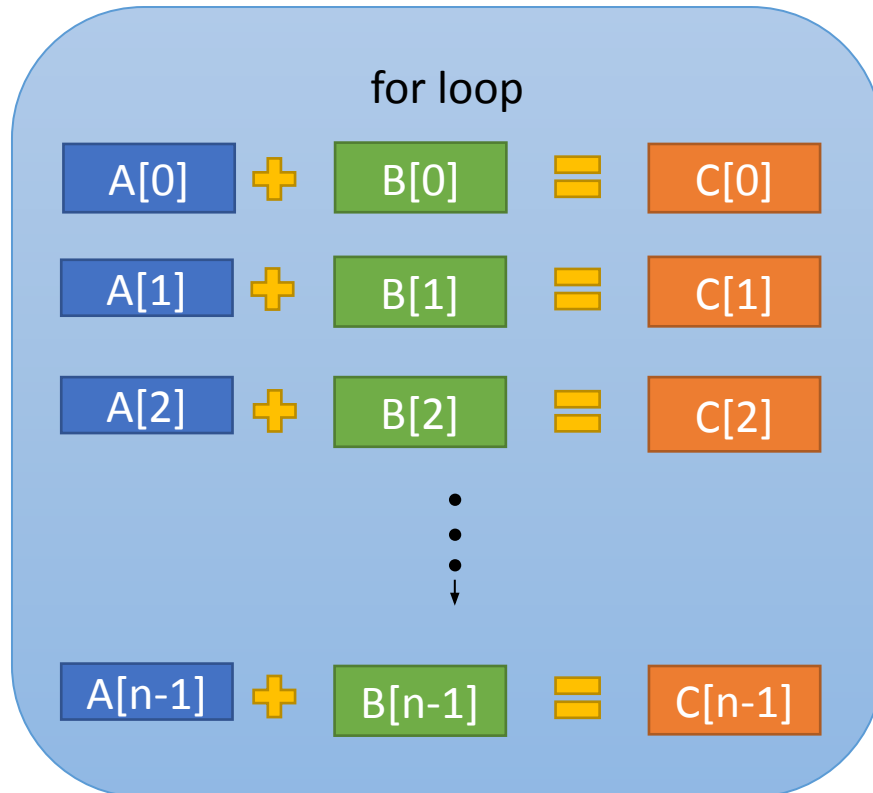
Introduction of GPU programming - vector addition - one thread

CUDA-enabled GPU



Vector Addition : vector_addition()

In one core as one thread




```
10  __global__ void vector_add(float *out, float *a, float *b, int n){  
11      for(int i=0;i<n;i++){  
12          out[i] = a[i] + b[i];  
13      }
```


Introduction to GPU

Introduction of GPU programming - vector addition - one streaming multiprocessor

Vector Addition : `vector_addition()`




- When the kernel is called, the **number of blocks** and the **number of threads** in each block is specified
`vector_add <<<1,256>>> (d_out, d_a, d_b, N)`
 - The block runs in a stream multiprocessor (SM)
 - The SM will have 256 thread.
 - Each thread runs in a core.
 - Each thread will have an unique id: `threadIdx.x`

`threadIdx.x` 0 1 2 254 255



Introduction of GPU programming - vector addition - one streaming multiprocessor

threadIdx.x 0 1 2 ... 255



blockDim.x 256

```
10  __global__ void vector_add(float *out, float *a, float *b, int n){
11      int index = threadIdx.x;
12      int stride = blockDim.x;
13
14      for(int i=index;i<n;i+=stride){
15          out[i] = a[i] + b[i];}
16  }
```

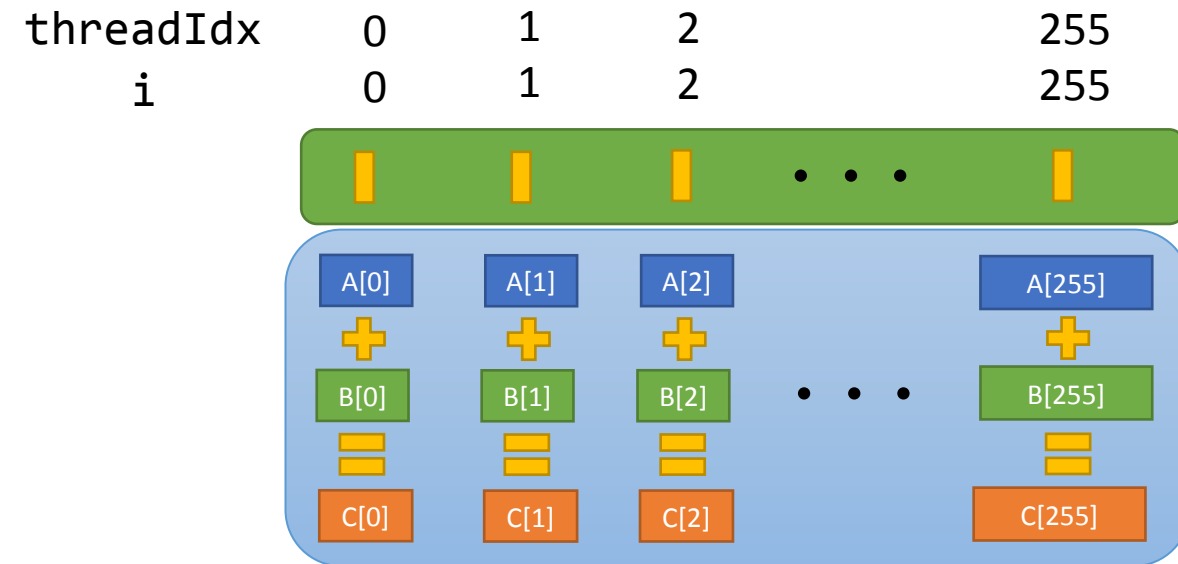
Introduction to GPU

Introduction of GPU programming - vector addition - one streaming multiprocessor

Vector Addition : vector_addition()

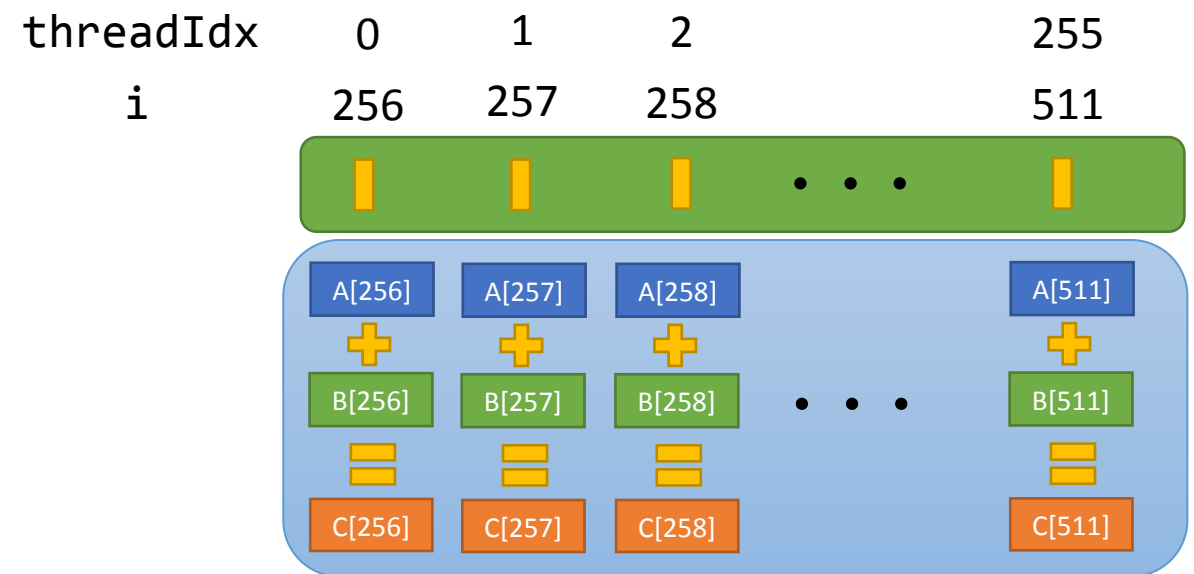
1st loop

$i = \text{threadIdx.x}$



2nd loop

$i = i + \text{blockDim.x}(=256)$



```
10  __global__ void vector_add(float *out, float *a, float *b, int n){
11      int index = threadIdx.x;
12      int stride = blockDim.x;
13
14      for(int i=index; i<n; i+=stride){
15          out[i] = a[i] + b[i];
16      }
```

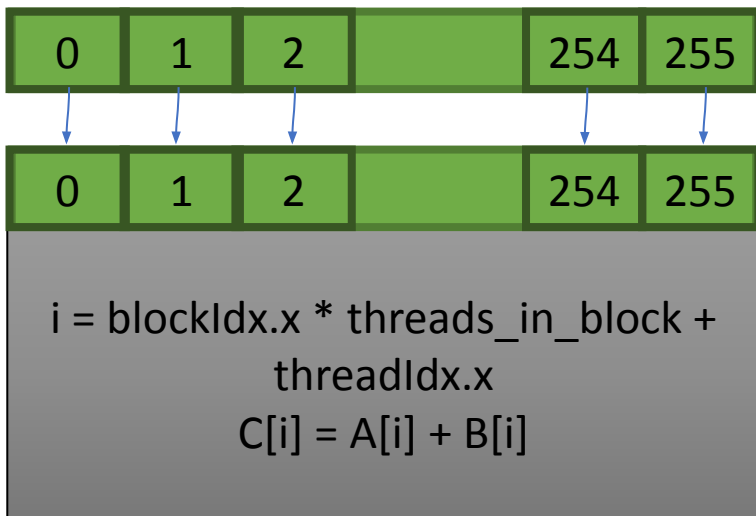
Introduction to GPU

Introduction of GPU programming - vector addition – GPU device


Vector Addition Example : `vector_addition()`

Number of blocks : N	<code>blockIdx.x</code> : 0, 1, 2, 3 ... N - 1
Number of threads in each block : 256	<code>threadIdx.x</code> : 0, 1, 2, 3 ... 255

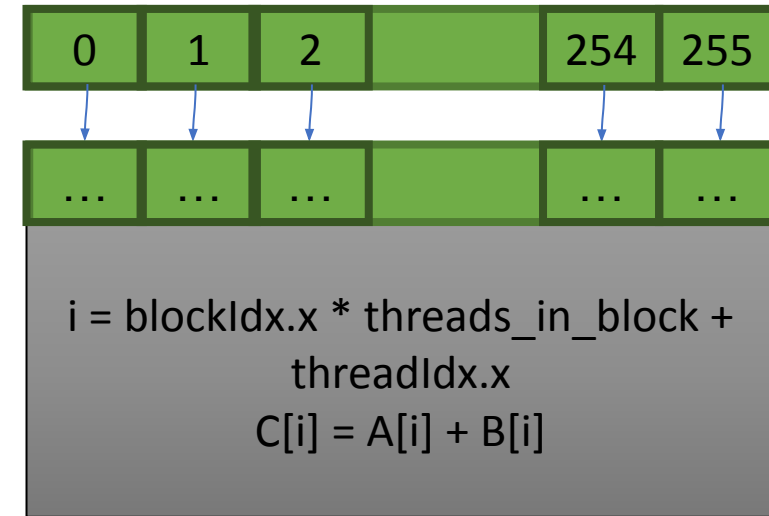
Block 0



Block 1



Block N-1



Introduction to GPU

Introduction to CUDA

CUDA programming Model – Heterogenous Computing

Host – CPU, Device - GPU

```
16 int main(){
17     float *a, *b, *out;
18     float *d_a, *d_b, *d_out;
19
20     a = (float*)malloc(sizeof(float) * N);
21     b = (float*)malloc(sizeof(float) * N);
22     out = (float*)malloc(sizeof(float) * N);
23
24     // Initialize array
25     for(int i = 0; i < N; i++){
26         a[i] = 1.0f;
27         b[i] = 2.0f;
28     }
29
30     // Allocate device memory for a
31     cudaMalloc((void**)&d_a, sizeof(float)*N);
32     cudaMalloc((void**)&d_b, sizeof(float)*N);
33     cudaMalloc((void**)&d_out, sizeof(float)*N);
34
35     // Transfer data from host to device memory
36     cudaMemcpy(d_a, a, sizeof(float)*N, cudaMemcpyHostToDevice);
37     cudaMemcpy(d_b, b, sizeof(float)*N, cudaMemcpyHostToDevice);
38
39     // Main function
40     int block_size = 256;
41     int grid_size = (N+block_size)/block_size;
42     vector_add<<<grid_size, block_size>>>>(d_out, d_a, d_b, N);
43
44     cudaMemcpy(out, d_out, sizeof(float)*N, cudaMemcpyDeviceToHost);
45
46     // Deallocate device memory
47     cudaFree(d_a);
48     cudaFree(d_b);
49     cudaFree(d_out);
50
51     // Deallocate host memory
52     free(a);
53     free(b);
54     free(out);
55
56 }
```

```
16 int main(){
17     float *a, *b, *out;
18     float *d_a, *d_b, *d_out;
19
20     a = (float*)malloc(sizeof(float) * N);
21     b = (float*)malloc(sizeof(float) * N);
22     out = (float*)malloc(sizeof(float) * N);
23
24     // Initialize array
25     for(int i = 0; i < N; i++){
26         a[i] = 1.0f;
27         b[i] = 2.0f;
28     }
29
30     // Allocate device memory for a
31     cudaMalloc((void**)&d_a, sizeof(float)*N);
32     cudaMalloc((void**)&d_b, sizeof(float)*N);
33     cudaMalloc((void**)&d_out, sizeof(float)*N);
34
35     // Transfer data from host to device memory
36     cudaMemcpy(d_a, a, sizeof(float)*N, cudaMemcpyHostToDevice);
37     cudaMemcpy(d_b, b, sizeof(float)*N, cudaMemcpyHostToDevice);
38
39     // Main function
40     int block_size = 256;
41     int grid_size = (N+block_size)/block_size;
42     vector_add<<<grid_size, block_size>>>>(d_out, d_a, d_b, N);
43
44     cudaMemcpy(out, d_out, sizeof(float)*N, cudaMemcpyDeviceToHost);
45
46     // Deallocate device memory
47     cudaFree(d_a);
48     cudaFree(d_b);
49     cudaFree(d_out);
50
51     // Deallocate host memory
52     free(a);
53     free(b);
54     free(out);
55
56 }
```

Memory Allocation
in Host

```
--
30 // Allocate device memory for a
31 cudaMalloc((void**)&d_a, sizeof(float)*N);
32 cudaMalloc((void**)&d_b, sizeof(float)*N);
33 cudaMalloc((void**)&d_out, sizeof(float)*N);
34
```

Memory Allocation
in Device

```
35 // Transfer data from host to device memory
36 cudaMemcpy(d_a, a, sizeof(float)*N, cudaMemcpyHostToDevice);
37 cudaMemcpy(d_b, b, sizeof(float)*N, cudaMemcpyHostToDevice);
--
```

Data transfer from
Host to Device

```
--
46 // Deallocate device memory
47 cudaFree(d_a);
48 cudaFree(d_b);
49 cudaFree(d_out);
50
51 // Deallocate host memory
52 free(a);
53 free(b);
54 free(out);
55
56
```

Deallocation of
Memory

```
cudaMemcpy(out, d_out, sizeof(float)*N, cudaMemcpyDeviceToHost);
```

Data transfer from
Device to Host

```
--
39 // Main function
40 int block_size = 256;
41 int grid_size = (N+block_size)/block_size;
42 vector_add<<<grid_size, block_size>>>>(d_out, d_a, d_b, N);
--
```

Computation in
Device

Introduction to GPU

GPU programming syntax

- Functions that run on GPU are usually enclosed in “<<< >>>”.
- It has extension ‘.cu’.
- Compiled using nvcc compiler driver.

Introduction to GPU

GPU programming syntax

The host device CPU manages memory

- It allocates the memory in the GPU
 - `cudaMalloc (void **pointer, size_t nbytes)`
 - `cudaMemset (void *point, int value, size_t count)`
- Data is transferred from host memory to device memory
 - `cudaMemcpy(device_variable, host_variable, size_of_variable, CudaMemcpyHostToDevice)`
- After the kernel execution and data is transferred from device to host memory
 - `cudaFree (void *pointer)`
- Finally, the memory is deallocated
 - `cudaMemcpy(host_variable, device_variable, size_of_variable, CudaMemcpyDevicetoHost)`

Introduction to GPU

GPU programming syntax

EXERCISE 1 : Vector Addition using GPU program

Source code:

https://github.com/sivasanarul/FEMwithGPU/tree/master/EX1_vector_addition

Introduction to GPU

Introduction of GPU

Code profiling – nvprof ./####

1 thread

```
==28571== Profiling application: ./vector_add_onethread
==28571== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max   Name
GPU activities:   94.28%   14.4019s        1  14.4019s  14.4019s  14.4019s  vector_add(float*, float*, float*, int)
                3.14%    479.91ms        2   239.95ms  239.17ms  240.74ms  [CUDA memcpy HtoD]
                2.58%    394.51ms        1   394.51ms  394.51ms  394.51ms  [CUDA memcpy DtoH]
```

1 block with 256 threads

```
==28894== Profiling application: ./vector_add_oneblock
==28894== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max   Name
GPU activities:   38.34%    478.81ms        2   239.40ms  239.19ms  239.62ms  [CUDA memcpy HtoD]
                30.86%    385.37ms        1   385.37ms  385.37ms  385.37ms  [CUDA memcpy DtoH]
                30.81%    384.72ms        1   384.72ms  384.72ms  384.72ms  vector_add(float*, float*, float*, int)
```


N blocks with all threads

```
==29270== Profiling application: ./vector_add_nblock
==29270== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max   Name
GPU activities:   55.40%    474.87ms        2   237.43ms  237.09ms  237.77ms  [CUDA memcpy HtoD]
                44.14%    378.35ms        1   378.35ms  378.35ms  378.35ms  [CUDA memcpy DtoH]
                 0.45%     3.8855ms        1    3.8855ms  3.8855ms  3.8855ms  vector_add(float*, float*, float*, int)
```

Introduction to GPU

Introduction of GPU programming

Vector Addition Example : `vector_addition()`



Introduction to GPU

Introduction of GPU

Code profiling

```
==29270== Profiling application: ./vector_add_nblock
==29270== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:  55.40%    474.87ms        2    237.43ms  237.09ms  237.77ms  [CUDA memcpy HtoD]
                  44.14%    378.35ms        1    378.35ms  378.35ms  378.35ms  [CUDA memcpy DtoH]
                  0.45%     3.8855ms        1     3.8855ms  3.8855ms  3.8855ms  vector_add(float*, float*, float*, int)
```

Expensive step is the memory transfer

To circumvent:

- Perform simultaneous data transfer and computation
- Some streaming multiprocessors perform computation, some perform data transfer